PS10 Finley

Caden Finley

April 2023

1 Workflow

library(tidyverse) library(tidymodels) library(magrittr) library(modelsummary) library(rpart) library(e1071) library(kknn) library(nnet) library(kernlab) library(readr)

set.seed(100)

income ¡- read$_csv$("$http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data$", $col_names = FALSE$)$names(income) <- -c$("age","workclass","fnlwgt","edu

Clean up the data Drop unnecessary columns income Make sure continuous variables are formatted as numeric income Make sure discrete variables are formatted as factors income Combine levels of factor variables that currently have too many levels income Advanced = c("Masters","Doctorate","Prof-school"), Bachelors = c("Bachelors"), SomeCollege = c("Some-college","Assoc-acdm","Assoc-voc"), HSgrad = c("HS-grad","12th"), HSdrop = c("11th","9th","7th-8th","1st-4th","10th","5th-6th","Preschool")), marital.status = fct$_collapse(marital.status, Married$ = c$("Married-civ-spouse","Married-spouse-absent","Married-AF-spouse"), Divorced = c$("Divorced","Separated"), Widowed = c$("Widowed"), NeverMarried = c$("Never-married")), race = fct$_collapse(race, White = c$("White"), Black = c$("Black"), Asian = c$("Asian-Pac-Islander"), Other = c$("Other","Amer-Indian-Eskimo")), workclass = fct$_collapse(workclass, Private = c$("Private"), SelfEmp = c$("Self-emp-not-inc","Self-emp-inc"), Gov = c$("Federal-gov","Local-gov","State-gov"), Other = c$("Without-pay","Never-worked","?")), occupation = fct$_collapse(occupation, BlueCollar$ = c("?","Craft-repair","Farming-fishing","Handlers-cleaners","Machine-op-inspct","Transport-moving"), WhiteCollar = c$("Adm-clerical","Exec-managerial","Prof-specialty","Sales","Tech-support"), Services = c$("Armed-Forces","Other-service","Priv-house-serv","Protective-serv")))

tidymodels time! income$_split <- -initial_split(income, prop = 0.8)income_train <- -training(income_split)income_test <- -testing(income_split)$

logistic regression print('Starting LOGIT') set up the task and the engine tune$_logit_spec <- -logistic_reg(penalty = tune(), tuningparametermixture = 1$1 = $lasso, 0 = ridge)$set$_engine$("glmnet")set$_mode$("classification")

define a grid over which to try different values of the regularization parameter lambda lambda$_grid <- -grid_regular(penalty(), levels = 50)$

3-fold cross-validation rec$_folds <- -vfold_cv(income_train, v = 3)$

Workflow $\text{rec}_wf < -workflow()add_model(tune_logit_spec)add_formula(high.earner\ education + marital.status + race + workclass + occupation + relationship + sex + age + capital.gain + capital.loss + hours)$

Tuning results $\text{rec}_res < -rec_wftune_grid(resamples = rec_folds, grid = lambda_grid)$

what is the best value of lambda? $\text{top}_acc < -show_best(rec_res, metric = "accuracy")best_acc < -select_best(rec_res, metric = "accuracy")final_logit_lasso < -finalize_workflow(rec_wf, best_acc)$

print('********** LOGISTIC REGRESSION **********') $\text{logit}_test < -last_fit(final_logit_lasso, income_split)collect_metrics()$

$\text{logit}_testtop_acc$

combine results into a nice tibble (for later use) $\text{logit}_ans < -top_acclogit_ansmutate(alg = "logit")$

tree model print('Starting TREE') set up the task and the engine $\text{tune}_tree_spec < -decision_tree(min_n = tune(), tuningparametertree_depth = tune(), tuningparametercost_complexity = tune(), tuningparameter)set_engine("rpart")set_mode("classification")$

define a set over which to try different values of tuning parameters $\text{tree}_grid < -grid_latin_hypercube(min_n(), tree_depth(), cost_complexity(), size = 20)$

3-fold cross-validation $\text{tree}_folds < -vfold_cv(income_train, v = 3)$

Workflow $\text{tree}_wf < -workflow()add_model(tune_tree_spec)add_formula(high.earner\ education + marital.status + race + workclass + occupation + relationship + sex + age + capital.gain + capital.loss + hours)$

Tuning results $\text{tree}_res < -tree_wftune_grid(resamples = tree_folds, grid = tree_grid)$

what are the best values of the tuning parameters? $\text{top}_acc_tree < -show_best(tree_res, metric = "accuracy")best_acc_tree < -select_best(tree_res, metric = "accuracy")final_tree < -finalize_workflow(tree_wf, best_acc_tree)$

print('********** DECISION TREE **********') $\text{tree}_test < -last_fit(final_tree, income_split)collect$

$\text{tree}_testtop_acc_tree$

combine results into a nice tibble (for later use) $\text{tree}_ans < -top_acc_treetree_ansmutate(alg = "tree")$

k-nearest neighbors print('Starting KNN') set up the task and the engine $\text{tune}_knn_spec < -nearest_neighbor(weight_func = tune(), tuningparameterneighbors = tune()tuningparameter)set_engine("kknn")set_mode("classification")$

define a set over which to try different values of tuning parameters $\text{knn}_grid < -grid_latin_hypercube(weight_func(), neighbors(), size = 20)$

3-fold cross-validation $\text{knn}_folds < -vfold_cv(income_train, v = 3)$

Workflow $\text{knn}_wf < -workflow()add_model(tune_knn_spec)add_formula(high.earner\ education + marital.status + race + workclass + occupation + relationship + sex + age + capital.gain + capital.loss + hours)$

Tuning results $\text{knn}_res < -knn_wftune_grid(resamples = knn_folds, grid = knn_grid)$

what are the best values of the tuning parameters? $\text{top}_acc_knn < -show_best(knn_res, metric = "accuracy")best_acc_knn < -select_best(knn_res, metric = "accuracy")final_knn < -finalize_workflow(knn_wf, best_acc_knn)$

print('********** K-NEAREST NEIGHBORS **********') $\text{knn}_t est <- last_f it(final_k nn, income_s plit) collect_m etrics()$

$\text{knn}_t est top_a cc_k nn$

combine results into a nice tibble (for later use) $\text{knn}_a ns <- top_a cc_k nn$

SVM

print('Starting SVM') set up the task and the engine $\text{tune}_s vm_s pec <- svm_p oly(degree = tune(), tuningparameter : degreeofpolynomialkernelscale = tune(), tuningparameter : scalingfactorforkernelcost = tune(), tuningparameter : costofviolationepsilon = 0.1epsilonforepsilon-insensitivelossfunction) set_e ngine("LiblineaR") set_m ode("classification")$

define a grid over which to try different values of tuning parameters $\text{svm}_g rid <- grid_t une(degree = seq(1, 5, by = 1), trydegreesfrom1to5scale = seq(0.1, 1, by = 0.1), tryscalingfactorsfrom0.1to1cost = 10^s eq(-3, 3, by = 1)trycostsfrom0.001to1000)$

Workflow $\text{svm}_w f <- workflow() add_m odel(tune_s vm_s pec) add_f ormula(high.earner\ education + marital.status + race + workclass + occupation + relationship + sex + age + capital.gain + capital.loss + hours)$

Tuning results $\text{svm}_r es <- svm_w f tune_g rid(resamples = rec_f olds, grid = svm_g rid)$

what is the best combination of tuning parameters? $\text{top}_a cc_s vm <- show_b est(svm_r es, metric = "accuracy") best_a cc_s vm <- select_b est(svm_r es, metric = "accuracy") final_s vm <- finalize_w orkflow(svm_w f, best_a cc_s vm) print('**********SUPPORTVECTORMACHINE(SVM)********** ') svm_t est <- last_f it(final_s vm, income_s plit) collect_m etrics()$

$\text{svm}_t est top_a cc_s vm$

combine results into a nice tibble (for later use) $\text{svm}_a ns <- top_a cc_s vmsvm_a nsmutate(alg = "svm")$

combine answers $\text{all}_a ns <- bind_r ows(logit_a ns, tree_a ns, nnet_a ns, knn_a ns, svm_a ns) datasummary_d f(all_a ns$

2 Explanation

I apologize for the messy code and lack of answers, I was not sure how to interpret some of what was required in the problem set. I had issues with the time constraint this time so I could not get it presented like I wanted to. Alongside that, getting this code to translate over in LaTex was giving me trouble, so again I apologize for the sloppy work.